

Siddharth Saminathan

+91 7299707403 | siddharthsaminathan99@gmail.com | [LinkedIn](#) | [emotenow.app](#) | Chennai, India

SKILLS

AI & Personalization

LLM APIs (OpenAI, Claude, Gemini) · RAG pipelines · embeddings · prompt design · multi-model routing · multi-agent orchestration · memory systems · conversational AI

Backend & Systems

Python · FastAPI · REST APIs · async pipelines · WebSockets · SSE · streaming · Redis (caching, queues) · PostgreSQL

Data

SQL · ETL pipelines · business intelligence · Power BI · Tableau · data validation · pipeline optimization

Infra & Deployment

Docker · Fly.io · AWS (basics) · logging · observability · performance optimization

EDUCATION

M.Sc. Statistics and Machine Learning, Linköping University
Sweden · 2020-2022

Thesis: Built an ML pipeline for colorectal cancer metastasis prediction from high-dimensional RNA sequencing data (~64,000 features, under 50 samples), demonstrating ability to extract signal from sparse, complex datasets.

Applied autoencoders for feature compression and representation learning, a technique transferable to building customer personalization intelligence layers.

Used PCA for dimensionality reduction and feature selection, enabling efficient model training.

Performed gene set enrichment analysis and statistical modelling to identify biologically relevant patterns.

PROFILE

AI-native product engineer with 6+ years of experience shipping consumer-facing AI systems from concept to real users. Co-founded and scaled a hyper-personalized AI companion to 450+ users with 13-15% weekly retention, achieving 50x cost reduction and 5x latency improvement through LLM API optimization. Proven ability to architect and deploy intelligence layers that drive measurable customer outcomes, retention, personalization, and brand experience. Entrepreneurial, CEO-facing communicator who thrives on building AI products that delight users and deliver business impact.

EXPERIENCE

AI Engineer

Omnex Systems
AquaPro AI

2025 - Present | Chennai, India

- Shipped an AI-native intelligence layer that personalizes quality workflows (DFMEA, PFMEA, Control Plans) by understanding product hierarchies and suggesting context-aware updates, reducing manual analysis effort for engineering teams.
- Architected a conversational BI system using LLM APIs and tool-calling: selects tools, routes data across internal services, and executes multi-step workflows end-to-end (generate, validate, update).
- Designed feedback-aware flows that retrieve similar past concerns, link them, update records, and suggest impactful actions, enabling a self-improving brand experience.
- Reduced response latency from ~12.5s to ~8s by restructuring the pipeline and consolidating model calls, improving real-time customer personalization.
- Implemented observability (structured logging, tracing, and metrics) to monitor system behavior and debug production issues.
- Improved output consistency and reduced failures with validation layers, retry handling, and fallback logic.
- Enabled the system to map inputs to the right structures and suggest updates based on existing data, ensuring accurate and relevant outputs.

B.Tech Computer Science & Engineering, SRM University
India · 2016-2020

COMPETENCIES

AI & Personalization

Backend & Systems

Data & Analytics

Infra & Deployment

Co-Founder, AI Product Engineer

Emote
emotenow.app

2024 - Present

- Shipped a production AI companion system that delivers hyper-personalized emotional support, ensuring users feel remembered across sessions, a pattern directly applicable to building a customer personalization intelligence layer for e-commerce.
- Designed a memory-driven personalization engine that tracks emotional context, relationships, and cause-effect patterns across 100s of conversations, improving relevance per user.
- Built the backend from scratch (FastAPI, Redis, PostgreSQL, hosted on Fly.io) with SSE streaming and WebSocket-based real-time pipelines, caching, and multi-model routing.
- Scaled from 0 to 450+ users, with 40+ power users and ~13-15% weekly cohort retention (pre-PMF).
- Reduced inference cost per conversation from ~\$1.20 to ~\$0.023 through model selection, caching, and prompt optimization.
- Improved latency from ~15s to ~3s through pipeline optimization and streaming.
- Built an internal AI product manager that monitors system behaviour, detects anomalies, and improves observability across the production AI stack.
- Built a Reddit bot to automate user acquisition; reduced outreach time from ~6-8 hours to ~30-40 minutes, enabling 50-100 hyper-personalized DMs per session, increasing activation from ~20% to ~75% and weekly retention from ~5% to ~13%.

Data Engineer

Akkodis AB
Ericsson Client

Nov 2022 - Nov 2024 | Sweden

- Designed and shipped an AI-native intelligence layer for business reporting, enabling conversational BI across dashboards used by stakeholders.
- Architected end-to-end ETL lifecycle for Power BI and Tableau dashboards, delivering real-time customer personalization insights.
- Drove data reliability and consistency during platform migration through automated quality audits, ensuring zero downtime for brand experience reporting.
- Owned redesign of data pipelines powering reporting systems, aligning data owners and stakeholders around a unified customer view.

Data Analyst

Positive Integers Pvt Ltd

Sep 2019 - May 2020 | Chennai, India

- Built Tableau dashboards for an NBFC to analyze credit risk and loan repayment behavior, enabling data-driven lending decisions.
- Shipped real-time Grafana dashboards for ID by SPI Cinemas to track the full order-to-table lifecycle, improving operational visibility.
- Owned SQL-based ETL workflows to extract, transform, and load data from operational source systems into reporting layers, ensuring data accuracy and availability. ---

KEY ACHIEVEMENTS

- Shipped Emote, an AI-native personalization app: grew to 450+ users with 40+ power users and 13-15% weekly retention through automated, personalized user acquisition. (<https://emotenow.app/>)
- Led enterprise data platform migration (Azure to PostgreSQL) at Akkodis for Ericsson; recognized for improving reporting reliability and delivery.
- --