

AI systems that get better with every interaction

I build production AI products where context, behavior, and workflow history compound over time – across users, teams, and business operations.

Siddharth Saminathan
AI Product Engineer · Builder · Chennai
[emotenow.app](#) · [GitHub](#)

01 - Evidence

Continuity drives return behavior

Emote is a live AI product I built from zero. The core bet was simple: people return when the system understands their context better over time. That bet held.

28.3%

WOW RETENTION AVG

Users returned week after week

April 2026

64.9%

ACTIVATION RATE

Nearly two-thirds reached meaningful usage

April 2026

492

TOTAL USERS

Built from zero, no paid acquisition

All-time

61

MONTHLY ACTIVE

Active, returning, engaged users

April 2026

27.4%

EFFECTIVE CHURN

74

NEW SIGNUPS

Measured and managed weekly

April 2026

Organic, referral-driven growth

April 2026

Wednesday

PEAK DAY

Consistent midweek engagement rhythm

50×

COST REDUCTION

Infrastructure optimized from prototype to production

All-time

Production metrics · May 2026

02 - Learning

What Emote taught me

01

Retention came from continuity, not novelty. Users returned when they felt the system understood prior context.

02

The product improved by studying behavior, not by adding random AI features. Every change earned its place in the workflow.

03

Latency and cost mattered because emotional products fail when they feel slow. Users notice sub-second delays in conversation.

04

The same principle applies to retail: every customer interaction should carry forward context across touchpoints.

05

People notice when a system remembers context correctly. That applies to a returning customer, a store associate, or a merchandising decision.

06

Shipping fast and measuring behavior beats shipping perfect and measuring opinions. Users showed us what mattered.

The important insight was not that people like AI. It was that people notice when a system remembers context correctly. That same principle applies to a returning customer, a store associate, a WhatsApp campaign, or a merchandising decision.

03 - Growth

Community-led acquisition without ad spend

I built a system that identified relevant community conversations, classified user intent, generated personalized outreach, tracked responses, and managed follow-ups - bringing in approximately 80% of Emote's user base without paid acquisition.

5,619

DMS SENT

Personalized, intent-matched outreach

25.9%

REPLY RATE

One in four responded - strong signal

~400

USERS ACQUIRED

From community outreach

80%

OF USER BASE

Community was the primary acquisition channel

1

Discover

2

Classify

3

Personalize

4

Send

5

Track

6

Follow Up

7

Learn

04 - Translation

Why this maps to Nicobar

The underlying systems are similar. Here is how my experience translates to what you are building.

CUSTOMER PERSONALIZATION

Context compounds across sessions

Emote proved that users stay when a system remembers who they are and what they care about. For Nicobar, the same principle becomes customer memory across email, WhatsApp, web, and store – every touchpoint carrying forward identity, preferences, and history.

STORE CLIENTELING

Interaction history guides better actions

The outreach CRM system classified intent and matched it to personalized communication. For Nicobar, this becomes associate-facing customer context: who is in front of them, what they love, what is in stock, what can ship from another store.

CONVERSATIONAL BI

Live questioning over static dashboards

My data engineering and dashboard work maps directly to Nicobar's vision of replacing static BI with a live conversational intelligence layer - merchandising, inventory, and campaign performance queried in plain language.

DESIGN WORKFLOWS

Remove friction so creative teams make more

My workflow automation experience translates to reducing creative friction: briefs, research, copy, asset variants, and content organization. The goal is the same - help a visual team move faster without losing what makes them Nicobar.

05 - Method

How I build

1

Start with the business question - define the right question before the first line of code

2

Ship a narrow prototype fast - get something in front of real users within days

3

Measure behavior, not opinions - let usage data guide what to build next

4

Improve until the system becomes part of the workflow - invisible infrastructure

I'm interested in Nicobar because the AI opportunity here is not to make the brand louder. It is to make every customer, store, and team interaction more contextual, more responsive, and more unmistakably Nicobar.

Siddharth Saminathan

AI Product Engineer

[emotenow.app](#) · [GitHub](#) · siddharthsaminathan99@gmail.com · +91 7299707403