

Siddharth Saminathan

+91 7299707403 | siddharthsaminathan99@gmail.com | [emotenow.app](#) | Chennai, India

Profile

AI Product and Systems Engineer with 4+ years of experience building and shipping production AI systems across data infrastructure, enterprise workflows, and consumer applications. Focus on taking systems from concept to real users, improving them through measurable outcomes such as latency, cost, and user retention. Experience spans multi-agent orchestration, real-time pipelines, and memory-driven AI systems, with a strong emphasis on reliability and iteration in production environments.

Competencies

Python

SQL

FastAPI

Redis (caching, queues)

PostgreSQL

REST APIs

async pipelines

WebSockets

SSE streaming

LLM APIs

RAG pipelines

embeddings

prompt design

multi-model routing

multi-agent orchestration

memory systems

Docker

Fly.io

AWS (basics)

logging

observability

performance optimization

ETL pipelines

Power BI

Tableau

data validation

pipeline optimization

--

Experience

AI Engineer

Omnex Systems (AquaPro AI), Chennai, India

2025 - Present

- Built a production AI system enabling AI-assisted quality and control workflows (DFMEA, PFMEA, Control Plans, Process Flows, 5W2H, 5 Whys, 8D), improving consistency and reducing manual analysis effort for engineering teams.
- Designed an orchestration layer using tool-calling: selects tools, routes data across internal services, and executes multi-step workflows end-to-end (generate, validate, update).
- Enabled the system to understand production hierarchies and categories, map inputs to the right structures, and suggest updates based on existing data.
- Implemented feedback-aware flows: retrieves similar past concerns, links them, updates records, and suggests impactful actions to users.
- Reduced response latency from ~12.5s to ~8s by restructuring the pipeline and consolidating model calls.
- Improved output consistency and reduced failures with validation layers, retry handling, and fallback logic.
- Implemented observability (structured logging, tracing, and metrics) to monitor system behavior and debug production issues.

Co-Founder, AI Product Engineer

Emote (emotenow.app)

2024 - Present

- Built a production AI companion system for hyper-personalized emotional support, users feel remembered and do not start over each session.
- Designed a memory system that tracks emotional context, relationships, and cause-effect patterns across 100s of conversations, improving relevance per user.
- Built the backend from scratch (FastAPI, Redis, PostgreSQL, hosted on Fly.io) with SSE streaming and WebSocket-based real-time pipelines, caching, and multi-model routing.
- Scaled from 0 to 450+ users, with 40+ power users and ~13-15% weekly cohort retention (pre-PMF).
- Reduced inference cost per conversation from ~\$1.20 to ~\$0.023 through model selection, caching, and prompt optimization.
- Improved latency from ~15s to ~3s through pipeline optimization and streaming.
- Built an internal AI product manager that monitors system behaviour, detects anomalies, and improves observability across the production AI stack.
- Built a Reddit bot to automate user acquisition; reduced outreach time from ~6-8 hours to ~30-40 minutes, enabling 50-100 hyper-personalized DMs per session, increasing activation from ~20% to ~75% and weekly retention from ~5% to ~13%.

Data Engineer

Akkodis AB (Ericsson Client), Sweden

Nov 2022 - Nov 2024

- Led migration of enterprise data platform from Azure to on-prem PostgreSQL cloud environment.
- Owned redesign of data pipelines powering reporting systems across data owners and stakeholders.
- Built and designed end-to-end ETL lifecycle for dashboards (Power BI, Tableau) used by business stakeholders.
- Ensured data reliability and consistency during large-scale migration through data quality audits.

Data Analyst

Positive Integers Pvt Ltd, Chennai, India

Sep 2019 - May 2020

- Built Tableau dashboards for an NBFC to analyze credit risk and loan repayment behaviour.
- Built real-time Grafana dashboards for ID by SPI Cinemas to track full order-to-table lifecycle.
- Built SQL-based ETL workflows to extract, transform, and load data from operational source systems into reporting layers.

Key Achievements

Built Emote from zero: grew to 450+ users with 40+ power users, 13-15% weekly retention with zero paid acquisition. (<https://emotenow.app/>)

Led enterprise data platform migration (Azure to PostgreSQL) at Akkodis for Ericsson; recognised for improving reporting reliability and delivery.

Education

M.Sc. Statistics and Machine Learning, Linköping University 2020-2022

Sweden Thesis: ML pipeline for colorectal cancer metastasis prediction from high-dimensional RNA sequencing data (~64,000 features, under 50 samples). Applied autoencoders for feature compression and representation learning. Used PCA for dimensionality reduction and feature selection. Performed gene set enrichment analysis and statistical modelling to identify biologically relevant patterns.

B.Tech Computer Science & Engineering, SRM University 2016-2020
India

Skills

Programming & Data	Python · SQL
Backend & Systems	FastAPI · Redis (caching, queues) · PostgreSQL · REST APIs · async pipelines · WebSockets · SSE streaming
AI / ML Systems	LLM APIs · RAG pipelines · embeddings · prompt design · multi-model routing · multi-agent orchestration · memory systems
Infra & Deployment	Docker · Fly.io · AWS (basics) · logging · observability · performance optimization
Data	ETL pipelines · Power BI · Tableau · data validation · pipeline optimization · --